

CSEE5590 Big Data Programming

In Class Programming –14 Report
(Jongkook Son)

Project Overview:

MLlib is Apache Spark's scalable machine learning library, with APIs in Java, Scala, Python, and R.

Requirements/Task(s):

In class exercise:

1. Classification:

Algorithms

a. Naïve Bayes b. Decision Tree c. Random Forest

2. Clustering:

Algorithm

a. KMeans

3. Regression:

Algorithm

a. Linear Regression b. Logistic Regression

What I learned in ICP:

I learned How to implement some ML algorithm using spark Mlib. Spark MLlib library provide scalable machine learning functionality which can leverage the Spark framework. ML offers the ability to make prediction based upon given datasets, and improves based upon the size of the dataset, which offers a nice synergy based upon the ability for Spark to handle very large amounts of data set. It was very meaningful to use spark for ML, which was very new to me.

Part 1 Classification

```
[ ] !pip install pyspark
```

Requirement already satisfied: pyspark in /usr/local/lib/python3.7/dist-packages (3.1.1)
Requirement already satisfied: py4j==0.10.9 in /usr/local/lib/python3.7/dist-packages (from pyspark) (0.10.9)

```
▶ from pyspark.ml.feature import VectorAssembler  
from pyspark.ml.classification import NaiveBayes  
from pyspark.ml.evaluation import MulticlassClassificationEvaluator  
#import numpy  
# Load training data  
from pyspark.ml.linalg import SparseVector  
# from pyspark.python.pyspark.shell import spark  
from pyspark.sql import SparkSession  
import os
```

```
[ ] os.environ["HADOOP_HOME"] = "C:/winutils"
```

```
[ ] # Creating spark session  
spark = SparkSession.builder.appName("ICP7").getOrCreate()  
spark.sparkContext.setLogLevel("ERROR")
```

First import libraries and create spark session

```
[ ] data = data.select("age", "capital-gain", "education-num", "hours-per-week", "income")  
data.show()
```

```
+---+-----+-----+-----+-----+  
|age|capital-gain|education-num|hours-per-week|income|  
+---+-----+-----+-----+-----+  
| 39|         2174|          13|          40|<=50K|  
| 50|           0|          13|          13|<=50K|  
| 38|           0|           9|          40|<=50K|  
| 53|           0|           7|          40|<=50K|  
| 28|           0|          13|          40|<=50K|  
| 37|           0|          14|          40|<=50K|  
| 49|           0|           5|          16|<=50K|  
| 52|           0|           9|          45|>50K|  
| 31|        14084|          14|          50|>50K|  
| 42|        5178|          13|          40|>50K|  
| 37|           0|          10|          80|>50K|  
| 30|           0|          13|          40|>50K|  
| 23|           0|          13|          30|<=50K|  
| 32|           0|          12|          50|<=50K|  
| 40|           0|          11|          40|>50K|  
| 34|           0|           4|          45|<=50K|  
| 25|           0|           9|          35|<=50K|  
| 32|           0|           9|          40|<=50K|  
| 38|           0|           7|          50|<=50K|  
| 43|           0|          14|          45|>50K|  
+---+-----+-----+-----+-----+  
only showing top 20 rows
```

Select numeric Features from data

Now, to train a model need to have features for training which I have already chosen, and a target that I wanted to predict.

For the target, I used income feature.

However, we need it as a numeric, so I used the stringIndexer to do feature scaling as shown below.

```
[ ] from pyspark.ml.feature import StringIndexer
    indexer = StringIndexer(inputCol="income", outputCol="label")
    data_converted = indexer.fit(data).transform(data)
```

```
data_converted.show()
```

age	capital-gain	education-num	hours-per-week	income	label
39	2174	13	40	<=50K	0.0
50	0	13	13	<=50K	0.0
38	0	9	40	<=50K	0.0
53	0	7	40	<=50K	0.0
28	0	13	40	<=50K	0.0
37	0	14	40	<=50K	0.0
49	0	5	16	<=50K	0.0
52	0	9	45	>50K	1.0
31	14084	14	50	>50K	1.0
42	5178	13	40	>50K	1.0
37	0	10	80	>50K	1.0
30	0	13	40	>50K	1.0
23	0	13	30	<=50K	0.0
32	0	12	50	<=50K	0.0
40	0	11	40	>50K	1.0
34	0	4	45	<=50K	0.0
25	0	9	35	<=50K	0.0
32	0	9	40	<=50K	0.0
38	0	7	50	<=50K	0.0
43	0	14	45	>50K	1.0

only showing top 20 rows

Then change all the features in to integer and used vector assembler

Finally I split the data 70% for train and 30% for test

Naïve Bayes



```
from pyspark.ml.classification import NaiveBayes
nbl = NaiveBayes()
```

```
# train the model
modell = nbl.fit(X_train)
```



```
# select example rows to display.
predictions = modell.transform(x_test)
predictions.show(5)
```



label	features	rawPrediction	probability prediction
0.0	[17.0,0.0,4.0,45.0]	[-125.35616342715...	[1.0,9.9608115215... 0.0
0.0	[17.0,0.0,5.0,15.0]	[-74.640681407274...	[1.0,1.5075237380... 0.0
0.0	[17.0,0.0,5.0,25.0]	[-92.611207198269...	[1.0,2.4944842373... 0.0
0.0	[17.0,0.0,5.0,35.0]	[-110.58173298926...	[1.0,4.1275977641... 0.0
0.0	[17.0,0.0,5.0,40.0]	[-119.56699588476...	[1.0,5.3095250096... 0.0

only showing top 5 rows



```
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

# compute accuracy on the test set
evaluator = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction",
                                              metricName="accuracy")
accuracy = evaluator.evaluate(predictions)
print("Test set accuracy = " + str(accuracy))
```

Test set accuracy = 0.7765089722675367

Decision Tree

```
[ ] from pyspark.ml.classification import DecisionTreeClassifier
DecisionTreeModel = DecisionTreeClassifier()
model2 = DecisionTreeModel.fit(x_test)
```

```
[ ] # Generate prediction from test dataset
dt = model2.transform(x_test)
```

```
[ ] # Evaluate the accuracy of the model
evaluator = MulticlassClassificationEvaluator()
accuracy = evaluator.evaluate(dt)
```

```
# Show model accuracy
print("Accuracy:", accuracy)
```

Accuracy: 0.8051962592361115

Random Forest

▼ Random Forest

```
[ ] from pyspark.ml.classification import RandomForestClassifier

    # Using the training set for the model training
    RandomForestModel = RandomForestClassifier()
    model3 = RandomForestModel.fit(x_test)
```

```
[ ] # Generate prediction from test dataset
    RF = model3.transform(x_test)
```

```
[ ] # Evaluate the accuracy of the model
    evaluator = MulticlassClassificationEvaluator()
    accuracy = evaluator.evaluate(RF)

    # Show model accuracy
    print("Accuracy:", accuracy)
```

```
Accuracy: 0.8069629875824761
```

Part 2 Clustering

```
[8] data = data.select("encounter_id", "patient_nbr", "admission_type_id", "discharge_disposition_id", "admission_source_id", "time_in_hospital", "num_lab_procedures", "nu

[9] # Create vector assembler for feature columns
assembler = VectorAssembler(inputCols=data.columns, outputCol="features")
data = assembler.transform(data)

[12] # Trains a k-means model.
kmeans = KMeans().setK(4).setSeed(1)
model = kmeans.fit(data)

# Make predictions
predictions = model.transform(data)

# Shows the result.
centers = model.clusterCenters()
print("Cluster Centers: ")
for center in centers:
    print(center)

Cluster Centers:
[2.52253555e+08 6.69269641e+07 1.75386082e+00 3.10166826e+00
5.45081030e+00 4.18875119e+00 4.36344137e+01 1.31282173e+00
1.86933746e+01 5.58198284e-01 2.9785100e-01 7.11058151e-01
7.99285033e+00]
[6.37113388e+07 2.36018150e+07 2.28603305e+00 4.73503742e+00
6.76842254e+00 4.60556419e+00 4.38627313e+01 1.39341901e+00
1.53232029e+01 1.86216636e-01 9.04674481e-02 5.78675328e-01
6.67843525e+00]
[1.54459882e+08 6.62242825e+07 2.02153280e+00 3.27142555e+00
4.99793172e+00 4.37988384e+00 4.23099589e+01 1.29689758e+00
1.61003258e+01 4.10454739e-01 2.23629409e-01 6.47967134e-01
7.60481655e+00]
[3.77976343e+08 9.43735415e+07 1.67479675e+00 2.96886774e+00
5.47392425e+00 4.14772953e+00 4.20316280e+01 1.35712869e+00
1.68027960e+01 4.75609756e-01 2.76422764e-01 6.34840373e-01
8.21108467e+00]
```

Part 3 Regression

Linear Regression

```
[12] from pyspark.ml.regression import LinearRegression
      model1 = LinearRegression(maxIter=10, regParam=0.3, elasticNetParam=0.8)
```

```
[14] # Fit the model
      model = model1.fit(data)
```

```
[15] # Print the coefficients and intercept for linear regression
      print("Coefficients: %s" % str(model.coef))
      print("Intercept: %s" % str(model.intercept))
```

```
➤ Coefficients: [0.3337891635819007, 0.5150505011624908]
      Intercept: 6.2559533571945725
```

```
🔍 # Summarize the model over the training set and print out some metrics
trainingSummary = model.summary
print("numIterations: %d" % trainingSummary.totalIterations)
print("objectiveHistory: %s" % str(trainingSummary.objectiveHistory))
trainingSummary.residuals.show()
print("RMSE: %f" % trainingSummary.rootMeanSquaredError)
print("r2: %f" % trainingSummary.r2)
```

```
➤ numIterations: 6
objectiveHistory: [0.5, 0.39634946019346834, 0.1536353726360829, 0.15116447772451408, 0.14653853943833373, 0.14653853141273573, 0.14653853141271792]
+-----+
| residuals |
+-----+
|-7.614301294335078|
|-7.614301294335078|
|-3.1364659885591237|
|-0.29946282271511393|
|-0.4024729229476236|
|-0.5846202873387085|
|-2.3517309624286753|
|-2.3517309624286753|
|-2.3517309624286753|
|-1.7091113364223958|
| 2.3548500461959776|
| 2.3548500461959776|
| 2.3548500461959776|
| 2.3548500461959776|
|-0.7969638019444448|
|-0.7969638019444448|
|-2.9162222883000624|
| 1.470500884750379|
| 3.588850441301048|
| 2.949104166452699|
+-----+
only showing top 20 rows

RMSE: 2.837581
r2: 0.780117
```

Logistic Regression

```
[17] # Load data and select feature and label columns
data = spark.read.format("csv").option("header", True).option("inferSchema", True).option("delimiter", ","),load("/content/drive/MyDrive/imports-85.csv")
```

```
import pandas as pd
pd.DataFrame(data.take(5), columns=data.columns).transpose()
```

```
[20] from pyspark.sql.functions import col, when

data = data.withColumn("label", when(col("num-of-doors") == "four", 1).otherwise(0)).select("length", "width", "height", "label")
```

```
[21] import pandas as pd
pd.DataFrame(data.take(5), columns=data.columns).transpose()
```

	0	1	2	3	4
length	168.8	168.8	171.2	176.6	176.6
width	64.1	64.1	65.5	66.2	66.4
height	48.8	48.8	52.4	54.3	54.3
label	0.0	0.0	0.0	1.0	1.0

```
[22] # Create vector assembler for feature columns
assembler = VectorAssembler(inputCols=data.columns[:2], outputCol="features")
data = assembler.transform(data)
```

```
[23] data = data.select("label", "features")
```

```
[24] from pyspark.ml.classification import LogisticRegression
model1 = LogisticRegression(maxIter=10, regParam=0.3, elasticNetParam=0.8)
```

```
[25] # Fit the model
model = model1.fit(data)
```

```
# Print the coefficients and intercept for linear regression
print("Coefficients: %s" % str(model.coefficients))
print("Intercept: %s" % str(model.intercept))
```

```
Coefficients: [0.0 0.0 0]
Intercept: 0.22533894187764542
```

```
[30] # Summarize the model over the training set and print out some metrics
trainingSummary = model.summary
```